


Link Prediction Data Challenge


Description and Key Points

The Task

Original Data



Challenge



Green : Known existing links (0)
Red : Known **non** existing links (1)
Grey : To predict status (0/1)

Besides the graph we also have information on the papers (title, authors, abstract)

Classification Task

- **training_set** <source,target,class> : pairs of papers and whether there is an edge between the two (1 means there is an edge between source and target)
- **node_information** : information on papers
- **testing_set** <source,target> : classify the pair of papers
- **example_simple_features** : example python code of the classification task

Example - Baseline

- Stem text : Porter Stemmer
- Simple similarities:
 - Title: overlapping terms
 - Authors: overlapping authors
 - Difference between publication years
- SVM classifier
 - Evaluate classifier on sample of data
 - k-fold validation

Evaluation

- F1 score: harmonic mean of precision and recall

- $F1 = 2 \frac{p*r}{p+r}$

- Precision : How many did I get correct $p = \frac{tp}{tp+fp}$

- Recall : How many of the desired class were correctly retrieved $r = \frac{tp}{tp+fn}$

Improving the Baseline

- Text based :
 - Cleaning:
 - Data analysis: Are all terms in the text needed?
 - Part of speech tagging for removing terms
 - TF-IDF based features
 - Topic modeling
- Graph Based:
 - Node properties : degree, pagerank
 - Clustering coefficient
 - k-core number
- Similarity metrics:
 - Jaccard Index
 - Cosine similarity

TF-IDF vectorization

- Sklearn TfidfVectorizer

```
vectorizer = TfidfVectorizer(stop_words="english")  
TFIDF_matrix = vectorizer.fit_transform(corpus)
```

- Corpus: a list of documents as strings
- TFIDF_matrix : Document X Terms (sparse) matrix
- Fit/transform:
 - Fit learns the dictionary terms and their importnacr
 - Transform computes TFIDF

TfidfVectorizer

- Parameters:
 - **Tokenizer** : Function to split strings into words. Override this if you want more than space based tokenization
 - **stop_words**: words to ignore
 - **ngram_range** : range of how many tokens should we use as “one term” (min,max)
 - **max/min_df** : **values in [0-1] range**; ignore words with frequency higher/lower than the specified
 - **vocabulary** : use only terms from this pre-computed list

Part of Speech Tagging

- Are all terms useful?

```
#assume tokens holds a set of words
tagged_tokens = nltk.pos_tag(tokens)
#tagged_tokens=[ (token, tag) ,...]
```

- (some) possible tags:
 - Verbs: VB,VBD,VBG,VCN,VBP,VBZ
 - Nouns: NN,NNP,NNPS,NNS
 - Adjectives: JJ,JJR,JJS
 - Other : *nltk.help.upenn_tagset()*

NetworkX

```
import networkx as nx
```

Import library

```
G = nx.DiGraph()
```

Create a new directed graph

```
G.add_edges_from([("A","B"), ("C","A")])
```

Add nodes and edges edges

```
print G.in_degree()  
print G.out_degree()
```


Print the in-degree and out-degree of the nodes

```
print G.neighbors("A")  
print G.neighbors("B")
```

Print the neighborhood nodes of A and B

Graph Properties

- Clustering Coefficient: $c_u = \frac{2T(u)}{\deg(u)(\deg(u)-1)}$
 - $T(u)$: the number of triangles u belongs to
- k-core number:
 - k-core : the maximal graph where all nodes have at least k neighbors
 - k-core number of node: the maximum k for which a node belongs to a k -core



Latent Semantic Analysis (LSA)

- Assume matrix A with TF-IDF values

```
U, S, V = np.linalg.svd(A)
```

- Intuition: The collection is produced by a collection of topics.
 - U : document to concept association. Each vector represents the topic weights for each document
 - S : A weight for each concept
 - V : term to concept association. Each vector represents the weight of a term to a topic
- We can project A directly to the concept space:

```
#keep only the most important topics  
M = np.dot(A, V[:k, :].transpose())
```

Similarity Metrics

- Jaccard Index (sets A, B) :

- $J(A, B) = \frac{|A \cap B|}{|A \cup B|}$

- Percentage of terms that overlap over possible terms

- Cosine similarity (vectors A, B):

- $similarity = \frac{A \cdot B}{|A||B|}$

- If we subtract the vector means then we get the Pearson Correlation

Further improvements

- Other classifiers
 - You can explore other classification algorithms from sklearn
- Optimizing hyper-parameters
 - Sklearn can automatically produce the scores of a classifier over a grid of possible values for the hyper-parameters (model_selection)

THE END

What else can you think?