

DNN(Deep Neural Network) for Natural Language Processing



Shujie Liu

Microsoft Research Asia

Natural Language Processing Group

Outline

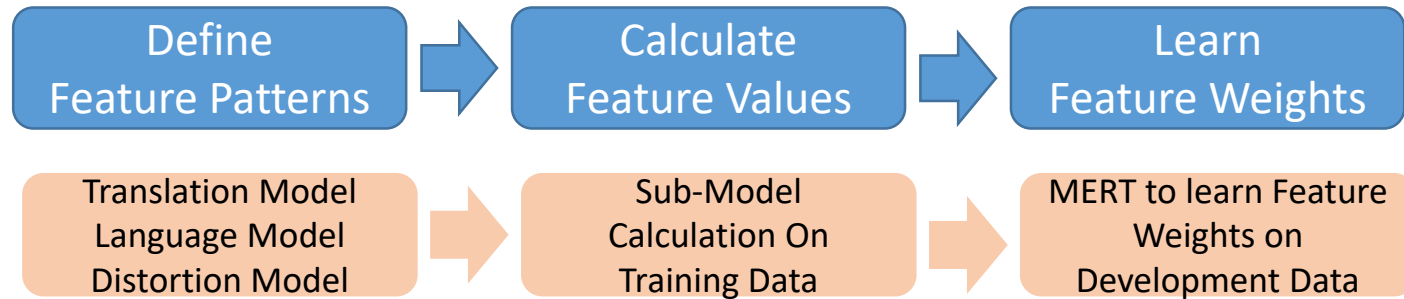
- Representation Learning
- Introduction to DNN
- DNN for Natural Language Processing

Outline

- Representation Learning
 - Why Learning Representation
 - How to Learn Representation
- Introduction to DNN
- DNN for Natural Language Processing

Why Learning Representation

- A popular way to build a ML model (crf, svm)



- Feature patterns/ Sub-models are defined by human knowledge.
- Feature engineering is important (most papers)
- Representation/Feature Learning
 - Learning information of the data that make it easier to extract useful information when building classifiers (Yoshua Bengio, et al., 2012)

How to Learn Representation

- PCA (Principal Components Analysis)
- Cluster Analysis
- Sparse Coding
- Deep Learning
-



Outline

- Representation Learning First
- Introduction to DNN
 - What is DNN
 - DNN vs. other models(perceptron, classifiers, etc.)
 - Why DNN now
 - Our understanding of DNN
- DNN for Natural Language Processing

Deep Neural Network

- **Deep Neural Network :**

- Involve multiple level neural networks
- Non-Linear Learner
- Automatically learn hierarchical representations from raw signals



DNN vs. Perceptron/SVM/MaxEnt

- Input Feature

- DNN consumes raw signals (or with minimal feature engineering)
- Others requires task specific, hand-crafted features

- Classifier Type

- DNN can handle highly non-linear space
- Others are essentially linear classifiers (SVM uses kernels to transform input space)

- What is learned

- DNN learns abstract representations (that can be shared with related tasks)
- Others learn feature weights (which is task-specific)


Why Now?



Big Data




Computing Power



DNN

- Training methods: SGD, Momentum, AdaGrad, AdaDelta, Adam
- Tricks: Parameter initialization, Early stopping, Dropout, Regularization, Active functions, Hyper-parameter optimization, Weight and Gradient normalization....

Back Propagation Training (by Example)



$$net_{h1} = w_1 * x_1 + w_2 * x_2 + b_1 * 1$$

$$out_{h1} = \frac{1}{1 + e^{-net_{h1}}}$$


$$net_{y1} = w_5 * net_{h1} + w_6 * net_{h2} + b_3 * 1$$

$$out_{y1} = \frac{1}{1 + e^{-net_{y1}}}$$

$$E = \sum \frac{1}{2} (target - output)^2$$

$$= \frac{1}{2} (target_{y1} - out_{y1})^2 + \frac{1}{2} (target_{y2} - out_{y2})^2$$

Back Propagation Training (by Example)



$$net_{h1} = w_1 * x_1 + w_2 * x_2 + b_1 * 1$$

$$out_{h1} = \frac{1}{1 + e^{-net_{h1}}}$$


$$net_{y1} = w_5 * net_{h1} + w_6 * net_{h2} + b_3 * 1$$

$$out_{y1} = \frac{1}{1 + e^{-net_{y1}}}$$

$$E = \sum \frac{1}{2} (target - output)^2$$

$$= \frac{1}{2} (target_{y1} - out_{y1})^2 + \frac{1}{2} (target_{y2} - out_{y2})^2$$

Back Propagation Training (by Example)



- Forward process:

$$net_{h1} = w_1 * x_1 + w_2 * x_2 + b_1 * 1$$


$$= 0.15 * 0.05 + 0.2 * 0.1 + 0.35 * 1 = 0.3775$$

$$out_{h1} = \frac{1}{1 + e^{-net_{h1}}} = \frac{1}{1 + e^{-0.3775}}$$

$$= 0.593269992$$

$$out_{h2} = 0.596884378$$

Back Propagation Training (by Example)



- Forward process:

$$net_{h1} = w_1 * x_1 + w_2 * x_2 + b_1 * 1$$

$$= 0.15 * 0.05 + 0.2 * 0.1 + 0.35 * 1 = 0.3775$$

$$out_{h1} = \frac{1}{1 + e^{-net_{h1}}} = \frac{1}{1 + e^{-0.3775}}$$

$$= 0.593269992$$

$$out_{h2} = 0.596884378$$

$$net_{y1} = w_5 * net_{h1} + w_6 * net_{h2} + b_3 * 1$$

$$= 0.4 * 0.593269992 + 0.45 * 0.596884378 + 0.6 * 1$$


$$= 1.105905967$$

$$out_{y1} = \frac{1}{1 + e^{-net_{y1}}} = \frac{1}{1 + e^{-1.105905967}}$$

$$= 0.75136507$$

$$out_{y2} = 0.772928465$$


Back Propagation Training (by Example)



- Calculate the error:

$$\begin{aligned}
 E &= \sum \frac{1}{2} (target - output)^2 \\
 &= \frac{1}{2} (target_{y_1} - out_{y_1})^2 \\
 &\quad + \frac{1}{2} (target_{y_2} - out_{y_2})^2 \\
 &= \frac{1}{2} (0.01 - 0.75136507)^2 \\
 &\quad + \frac{1}{2} (0.99 - 0.772928465)^2 \\
 &= 0.274811083 + 0.023560026 \\
 &= 0.298371109
 \end{aligned}$$

Back Propagation Training (by Example)



• Backward Process:


$$\frac{\partial E}{\partial w_5} = \frac{\partial E}{\partial out_{y1}} * \frac{\partial out_{y1}}{\partial net_{y1}} * \frac{\partial net_{y1}}{\partial w_5}$$

$$E = E_{y1} + E_{y2}$$


$$E_{y1} = \frac{1}{2} (target_{y1} - out_{y1})^2$$

$$out_{y1} = \frac{1}{1 + e^{-net_{y1}}}$$

$$net_{y1} = w_5 * net_{h1} + w_6 * net_{h2} + b_3 * 1$$



Back Propagation Training (by Example)




• Backward Process:

$$\frac{\partial E}{\partial w_5} = \frac{\partial E}{\partial out_{y1}} * \frac{\partial out_{y1}}{\partial net_{y1}} * \frac{\partial net_{y1}}{\partial w_5}$$


$$\frac{\partial E}{\partial out_{y1}} = -2 * \frac{1}{2} (target_{y1} - out_{y1})$$

$$\frac{\partial out_{y1}}{\partial net_{y1}} = out_{y1} (1 - out_{y1})$$

$$\frac{\partial net_{y1}}{\partial w_1} = out_{h1}$$



Back Propagation Training (by Example)




• Backward Process:

$$\frac{\partial E}{\partial w_5} = \frac{\partial E}{\partial out_{y1}} * \frac{\partial out_{y1}}{\partial net_{y1}} * \frac{\partial net_{y1}}{\partial w_5}$$


$$= -(target_{y1} - out_{y1})$$

$$* out_{y1}(1 - out_{y1})$$

$$* out_{h1}$$



Back Propagation Training (by Example)



- Backward Process:

$$\frac{\partial E}{\partial w_5} = \frac{\partial E}{\partial out_{y1}} * \frac{\partial out_{y1}}{\partial net_{y1}} * \frac{\partial net_{y1}}{\partial w_5}$$


$$= -(target_{y1} - out_{y1}) * out_{y1}(1 - out_{y1}) * out_{h1}$$

δ_{y1}
 out_{h1}

$$\delta_{y1} = (target_{y1} - out_{y1}) * out_{y1}(1 - out_{y1})$$

$$\frac{\partial E}{\partial w_5} = \delta_{y1} * out_{h1} = 0.082167041$$

Back Propagation Training (by Example)



• Backward Process:

$$\frac{\partial E}{\partial w_5} = \frac{\partial E}{\partial out_{y1}} * \frac{\partial out_{y1}}{\partial net_{y1}} * \frac{\partial net_{y1}}{\partial w_5}$$

$$= -(target_{y1} - out_{y1}) * out_{y1}(1 - out_{y1}) * out_{h1}$$

δ_{y1}
 out_{h1}

$$\delta_{y1} = (target_{y1} - out_{y1}) * out_{y1}(1 - out_{y1})$$


$$\frac{\partial E}{\partial w_5} = \delta_{y1} * out_{h1} = 0.082167041$$

• Update Process:

$$w_5^+ = w_5 - \alpha * \frac{\partial E}{\partial w_5} =$$

$$0.4 - 0.5 * 0.082167041 = 0.35891648$$

Back Propagation Training (by Example)




- Backward Process and Update Process for w_6 w_7 w_8 :

$$w_6^+ = 0.408666186$$

$$w_7^+ = 0.511301270$$

$$w_8^+ = 0.561370121$$

Back Propagation Training (by Example)



• Backward Process:

$$\frac{\partial E}{\partial w_1} = \frac{\partial E}{\partial out_{h1}} * \frac{\partial out_{h1}}{\partial net_{h1}} * \frac{\partial net_{h1}}{\partial w_1}$$

$$\frac{\partial E}{\partial out_{h1}} = \frac{\partial E_{y1}}{\partial out_{h1}} + \frac{\partial E_{y2}}{\partial out_{h1}}$$

$$\frac{\partial E_{y1}}{\partial out_{h1}} = \frac{\partial E_{y1}}{\partial net_{y1}} * \frac{\partial net_{y1}}{\partial out_{h1}}$$


$$\frac{\partial E_{y1}}{\partial net_{y1}} = 0.138498562$$

$$\frac{\partial net_{y1}}{\partial out_{h1}} = w_5$$

$$\begin{aligned} \frac{\partial E_{y1}}{\partial out_{h1}} &= 0.138498562 * 0.4 \\ &= 0.055399425 \end{aligned}$$

$$\frac{\partial E_{y2}}{\partial out_{h1}} = -0.019049119$$

Back Propagation Training (by Example)



• Backward Process:

$$\frac{\partial E}{\partial w_1} = \frac{\partial E}{\partial out_{h1}} * \frac{\partial out_{h1}}{\partial net_{h1}} * \frac{\partial net_{h1}}{\partial w_1}$$

$$\begin{aligned} \frac{\partial E}{\partial out_{h1}} &= \frac{\partial E_{y1}}{\partial out_{h1}} + \frac{\partial E_{y2}}{\partial out_{h1}} \\ &= 0.055399425 - 0.019049119 \\ &= 0.036350306 \end{aligned}$$


$$out_{h1} = \frac{1}{1 + e^{-net_{h1}}}$$

$$\frac{\partial out_{h1}}{\partial net_{h1}} = out_{h1}(1 - out_{h1}) = 0.241300709$$

$$net_{h1} = w_1 * x_1 + w_2 * x_2 + b_1 * 1$$

$$\frac{\partial net_{h1}}{\partial w_5} = x_1 = 0.05$$

Back Propagation Training (by Example)




• Backward Process:

$$\begin{aligned} \frac{\partial E}{\partial w_1} &= \frac{\partial E}{\partial out_{h1}} * \frac{\partial out_{h1}}{\partial net_{h1}} * \frac{\partial net_{h1}}{\partial w_1} \\ &= \left(\sum_i \frac{\partial E}{\partial out_{yi}} * \frac{\partial out_{yi}}{\partial net_{yi}} * \frac{\partial net_{yi}}{\partial w_1} \right) * \frac{\partial out_{h1}}{\partial net_{h1}} \\ &= \left(\sum_i \delta_{yi} * w_{hi} \right) * out_{h1} (1 - out_{h1}) * x_1 \\ &= \delta_{h1} * x_1 = 0.000438568 \end{aligned}$$

• Update Process:

$$\begin{aligned} w_1^+ &= w_1 - \alpha * \frac{\partial E}{\partial w_1} \\ &= 0.15 - 0.5 * 0.000438568 = 0.149780716 \end{aligned}$$


Back Propagation Training (by Example)



• Backward Process:

$$\begin{aligned} \frac{\partial E}{\partial w_1} &= \frac{\partial E}{\partial out_{h1}} * \frac{\partial out_{h1}}{\partial net_{h1}} * \frac{\partial net_{h1}}{\partial w_1} \\ &= \left(\sum_i \frac{\partial E}{\partial out_{yi}} * \frac{\partial out_{yi}}{\partial net_{yi}} * \frac{\partial net_{yi}}{\partial w_1} \right) * \frac{\partial out_{h1}}{\partial net_{h1}} \\ &= \left(\sum_i \delta_{yi} * w_{hi} \right) * out_{h1} (1 - out_{h1}) * x_1 \\ &= \delta_{h1} * x_1 = 0.000438568 \end{aligned}$$

Back Propagation Training (by Example)



- Backward Process and Update Process for w_2 w_3 w_4 :

$$w_2^+ = 0.19956143$$


$$w_3^+ = 0.24975114$$

$$w_4^+ = 0.29950229$$

- Compute the error using new parameters:

$$E = 0.298371109 \rightarrow 0.291027$$

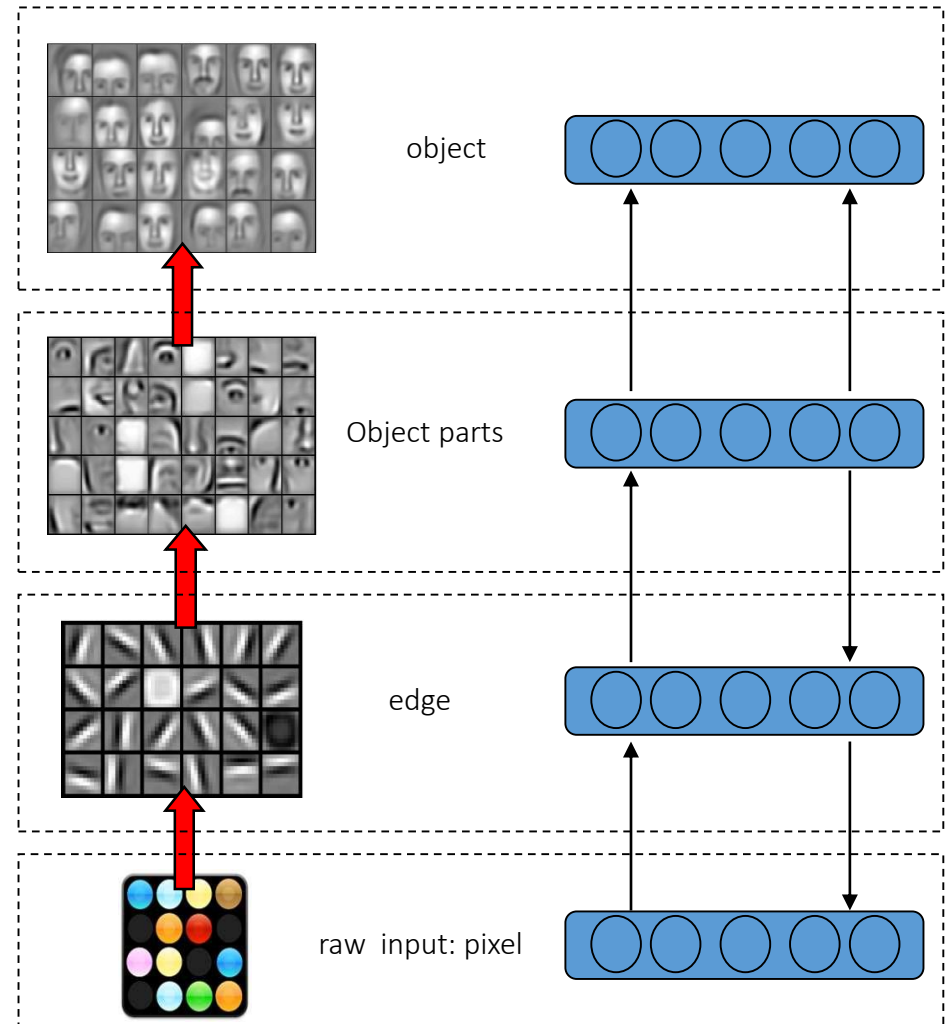

Back Propagation Training (the Rule)




α_{in} is the activation of the neuron input to the weight w

δ_{out} is the error of the neuron output from the weight w

DNN in Image Processing




CD-GMM-HMM for Speech



- GMM: Gaussian Mixture Model, used to model observation probability

CD-DNN-HMM for Speech



- GMM: Gaussian Mixture Model, used to model observation probability
- DNN is used to replace GMM
- Forced Alignment with CD-GMM-HMM is used to get training data for DNN
- Training data: tri-phone/signal-frame pairs

What DNN can do:

Representation/Feature Learning

- Finding multiple levels of representations
 - Dimensionality reduction: Denoising
 - Dimensionality ascension: Sparseness
 - Linear or none-linear mapping: (Kernel Function?)
- Minimal feature engineering
 - Representation(feature) varies for different tasks.
 - Transform 'factual knowledge into usable knowledge'
 - Learn good representations shared across multi tasks.

Advantage and Disadvantage of DNN

- Advantage
 - Usable features are learned **automatically**.
 - **Multiple levels** (coarse to fine) of representation.
- Disadvantage
 - Configuration and architecture is art: **too many hyper-parameters and variants**.
 - Features learned are **not understandable**.
 - Computational Complexity: **time consuming**

DNN in NLP


- NLP specific challenge:
 - Extremely high-dimensional space
 - Raw features (words) with strong discriminative power
 - Often highly structured output (Parsing, SMT)

Outline

- Representation Learning First
- Introduction to DNN
- DNN for Natural Language Processing
 - **DNN for Word Embedding**
 - DNN for Language Modeling
 - DNN for Machine Translation

Words into Vectors (Word Embedding)

- Words are embedded in a vector space



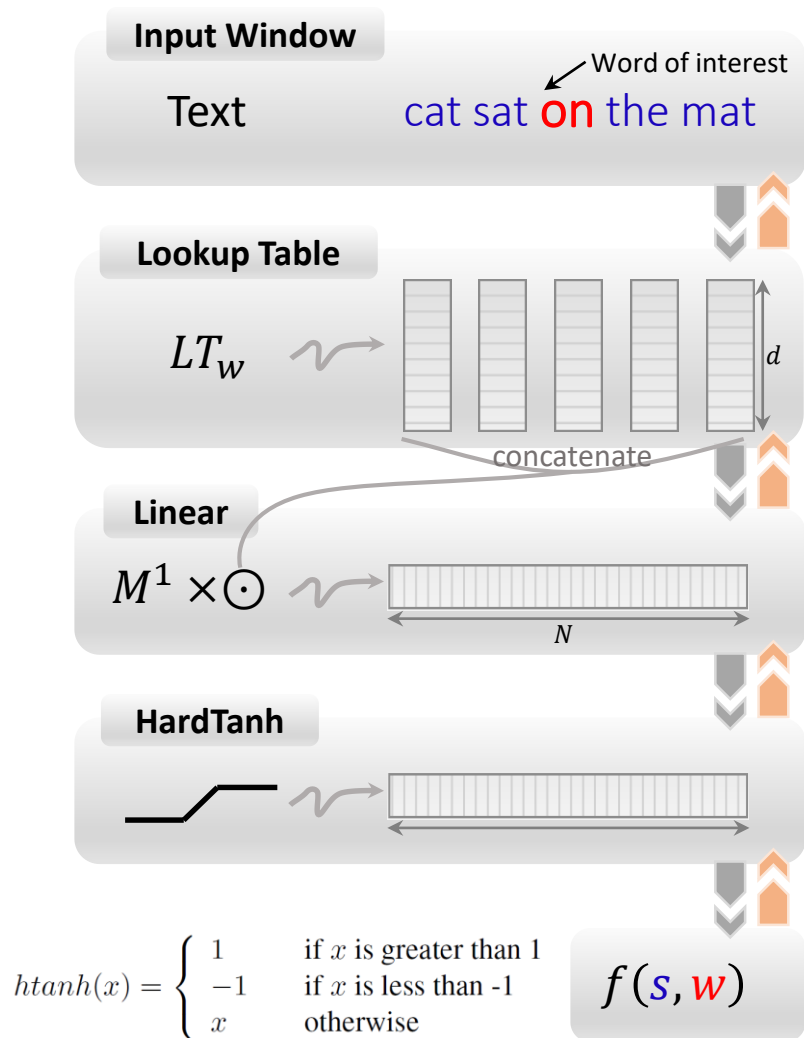
- Embeddings are **trained**
- Training Principle: "You shall know a word by the company it keeps" (J. R. Firth 1957)

government debt problems turning into banking crises as has happened in
saying that Europe needs unified banking regulation to replace the hodgepodge

↖ These words will represent *banking* ↗

You can vary whether you use local or large context to get a more syntactic or semantic clustering

JMLR11



- Ranking margin loss

$$\max(0, 1 - f(s, w_t^*) + f(s, w_t))$$

s : sentence window

w_t^* : true middle word in s

w_t : a random word to replace w_t^*

$f(s, w)$: network score for s and w


- Stochastic training and back propagation are used to train the parameters including the word embedding

CW08:Feed Forward NNLM



- Computational Complexity

$$Q = N \times D + N \times D \times H + H \times V$$

- The output is a probability
- Back propagation can also be used to update parameters




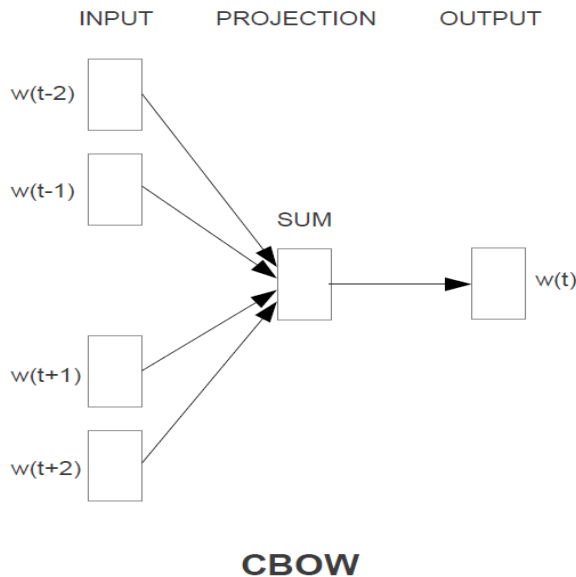
CW08 VS JMLR11





Google's Word2vec

- CBOW: Continuous Bag-of-Words
 - The order of words in the history does not influence the projection
- Computational Complexity



$$\cancel{Q = N \times D + N \times D \times H + H \times V}$$
$$Q = N \times D + D \times V$$





Linguistic Regularities in Continuous Space Word Representations (Mikolov, et al. 2013)

- Measuring Linguistic Regularity
 - Syntactic/Semetic Test



These representations are surprisingly good at capturing syntactic and semantic regularities in language, and that each relationship is characterized by a relation-specific vector offset.

Word Embedding Results (web 130G)

ipad
iphone 0.052583
psp 0.261444
ios 0.285222
xbox 0.2
zune 0.3
android 0.3
imac 0.3
ipod 0.3
tablet 0.3
wii 0.3
desktop 0.3
playstation
blackberry
app 0.4
pc 0.4
os 0.4
smartphone
firefox 0.4
laptop 0.4

zemin
qichen 0.104208
jintao 0.11451

metric tons **is smaller than** what
logo **seem smaller than** that one
are now **maybe smaller than** before
sign **is smaller than** the one next
were **much cheaper than**
, **cheaper than** in the States
and are **much cheaper than** organic fertilizers
sometimes **much cheaper than** the A shares

ceo
cfo 0.193586
founder 0.247986
chairman 0.254936
president 0.305918
owner 0.323968
publisher 0.328711
developer 0.330035
director 0.342627
analyst 0.34414
manager 0.365071
producer 0.376013
cto 0.381128
co-founder 0.38457
coo 0.400879
chief 0.40133
vp 0.429992

Bosses of
Company

cheaper
smaller 0.338234
stronger 0.338306
expensive 0.368454
smarter 0.374804
faster 0.380407
larger 0.401047
profitable
safer
pricier
costly
attractive 0.429141
inexpensive 0.431263
pricier 0.432898
affordable 0
usable 0.438529
cheap 0.448377

comparative
adjective

funny
silly 0.178232
scary 0.179388
weird 0.189012
boring 0.234093
sexy 0.237722
creepy 0.263487

stupid 0.285263
hilarious 0.291917
curious 0.315346
awkward 0.319954
bizarre 0.334614
ugly 0.340006


feeling
adjective

Not semantic similarity but syntactic similarity

HMM for Word Alignment

- Alignment Problem

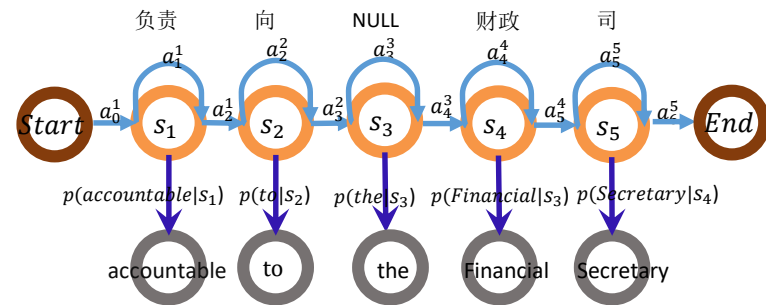
- Finding translation pairs in bitext sentences



- HMM for Word Alignment


- (Vogel et al., 1996)

$$p(a, e|f) = \prod_{i=1}^{|e|} p_{lex}(e_i|f_{a_i})p_d(a_i - a_{i-1})$$



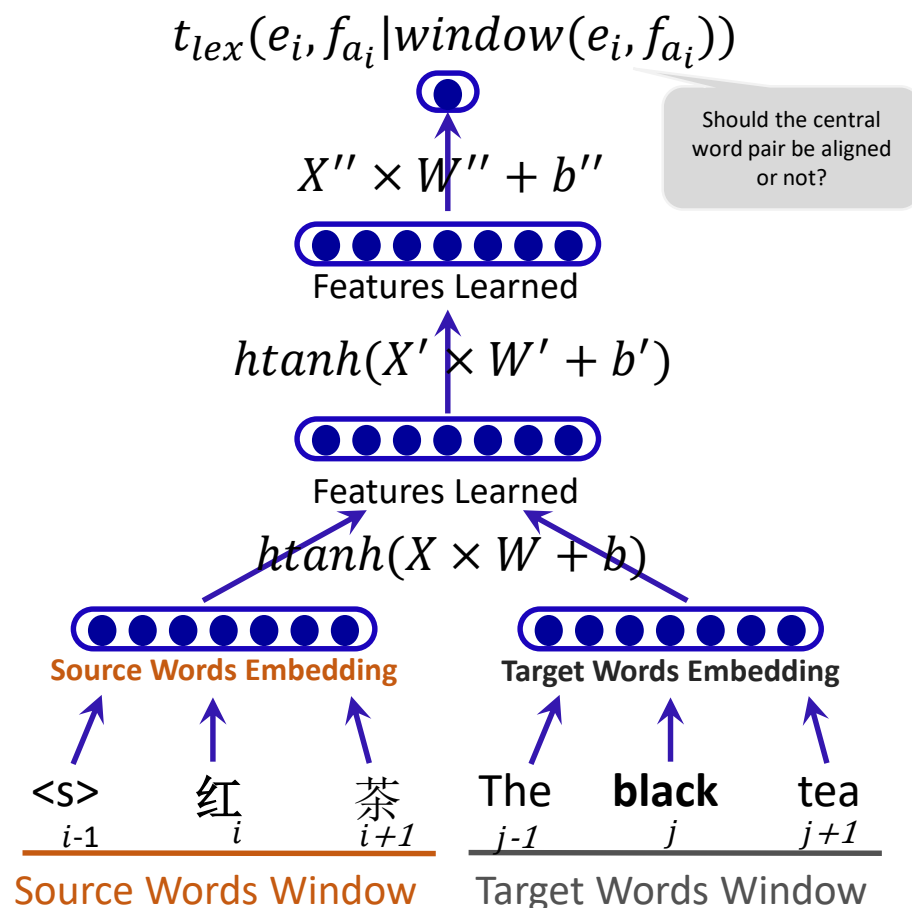
CD-DNN-HMM for Word Alignment

- Discover bilingual lexical similarity from mono and bi-text
- Improve language structure projections with rich context representation
- Cascaded training of neural latent variable models based on a initial alignment result



Model Observation Probability with DNN

- Source and target word embedding is initialized with word embedding trained with mono-lingual corpus
- Content information is used for training
- Embedding for both source and target words will be updated



Monolingual vs Bilingual Embedding

word	good	history	british	served	labs	zetian	laggards
LM	bad	tradition	russian	worked	networks	hongzhang	underperformers
	great	culture	japanese	lived	technologies	yaobang	transferees
	strong	practice	dutch	offered	innovations	keming	megabanks
	true	style	german	delivered	systems	xingzhi	mutuals
	easy	literature	canadian	produced	industries	ruihua	non-starters
WA	nice	historical	uk	offering	lab	hongzhang	underperformers
	great	historic	britain	serving	laboratories	qichao	illiterates
	best	developed	english	serve	laboratory	xueqin	transferees
	pretty	record	classic	delivering	exam	fuhuan	matriculants
	excellent	recording	england	worked	experiments	bingkun	megabanks

- “bad” is no longer in the nearest neighborhood of “good”, as they hold opposite semantic meaning
- Neighbors of proper nouns such as person names are relatively unchanged
- Rare words still remain their monolingual embeddings as they are modified a few times during bilingual training


Outline

- Representation Learning First
- Introduction to DNN
- DNN for Natural Language Processing
 - DNN for Word Embedding
 - **DNN for Language Modeling**
 - DNN for Machine Translation

Language Model

- What is Language Model?
 - Language Model is a model to tell whether a sequence of words is a sentence or not.
 - Probabilistic Language Modeling: Score the probability of a sentence generated.
 - $P(w_1, w_2, \dots, w_n) = p(w_1)p(w_2 | w_1) \dots p(w_n | w_1, \dots, w_{n-1})$
 - Too much parameters, independence assumption
 - Bigram: $P(w_1, w_2, \dots, w_n) = p(w_1)p(w_2 | w_1)p(w_3 | w_2) \dots p(w_n | w_{n-1})$
 - Trigram: $P(w_1, w_2, \dots, w_n) = p(w_1)p(w_2 | w_1)p(w_3 | w_2 w_1) \dots p(w_n | w_{n-1} w_{n-2})$
 - How to compute $p(w_2 | w_1)$ and $p(w_3 | w_2 w_1)$?
 - Maximum Likelihood Estimation (MLE) : $p(w_i | w_{i-1}) = \frac{\text{count}(w_{i-1}, w_i)}{\sum_{w^*} \text{count}(w_{i-1}, w^*)}$
 - if $\text{count}(w_{i-1}, w_i) = 0$?
 - Smoothing: add-1, gooden turing, Katz's back-off....

Feed Forward Language Model




- Discriminative LM
- Assign score for each word given context words with NN
- Produce word embedding (fixed length, real valued vector) which can be used by other NLP tasks
- Softmax Layer to generate the probability:

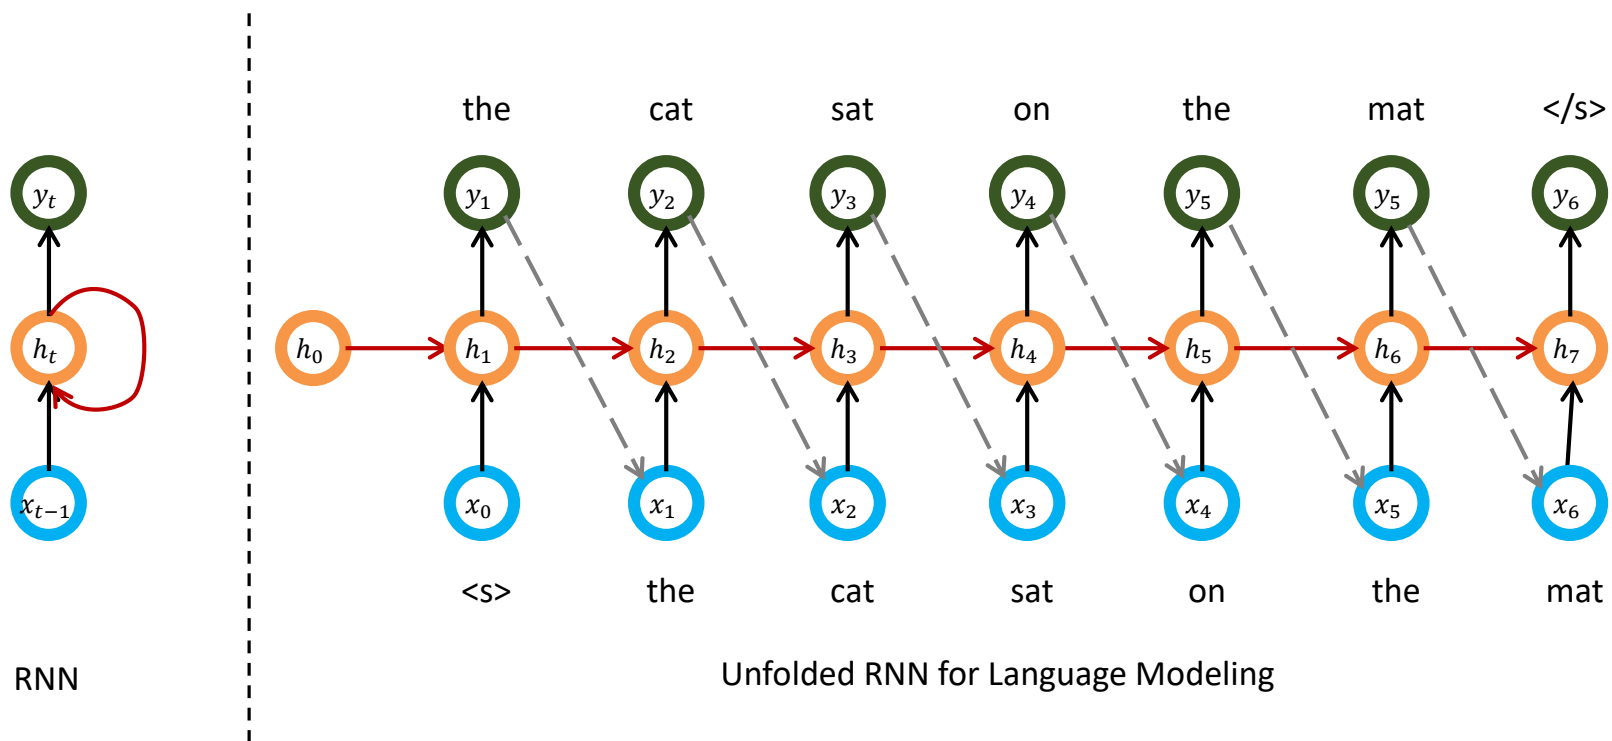
$$P(w_i) = \frac{e^{s_i}}{\sum_j e^{s_j}}$$

Recurrent Neural Network

- Inputs: History s_{t-1} at time $t - 1$ and input w_t at time t
- Output: History s_t at time t and next input y_t at time $t+1$




Unfolded RNN



Backward Propagation Through Time

$t = 1$


$$\delta_{in,1} = \frac{\partial L}{\partial h_{out,1}} \times \frac{\partial h_{out,1}}{\partial h_{in,1}}$$



Backward Propagation Through Time


t = 1

$$\delta_{in,1} = \frac{\partial L}{\partial h_{out,1}} \times \frac{\partial h_{out,1}}{\partial h_{in,1}}$$




t = 2

$$\begin{aligned} \delta_{in,1} &= \frac{\partial L}{\partial h_{out,2}} \times \frac{\partial h_{out,2}}{\partial h_{in,2}} \times \frac{\partial h_{in,2}}{\partial h_{out,1}} \times \frac{\partial h_{out,1}}{\partial h_{in,1}} \\ &= \delta_{out,2} \times \frac{\partial h_{out,2}}{\partial h_{in,2}} \times \frac{\partial h_{in,2}}{\partial h_{out,1}} \times \frac{\partial h_{out,1}}{\partial h_{in,1}} \\ &= \delta_{in,2} \times \frac{\partial h_{in,2}}{\partial h_{out,1}} \times \frac{\partial h_{out,1}}{\partial h_{in,1}} = \delta_{out,1} \times \frac{\partial h_{out,1}}{\partial h_{in,1}} \end{aligned}$$




Backward Propagation Through Time

$$\delta_{in,s} = \begin{cases} \frac{\partial L}{\partial h_{out,s}} \times \frac{\partial h_{out,s}}{\partial h_{in,s}} & \text{if } s = t \\ \delta_{in,s+1} \times \frac{\partial h_{in,s+1}}{\partial h_{out,s}} \times \frac{\partial h_{out,s}}{\partial h_{in,s}} & \text{otherwise} \end{cases}$$




Vanishing Gradient Problem

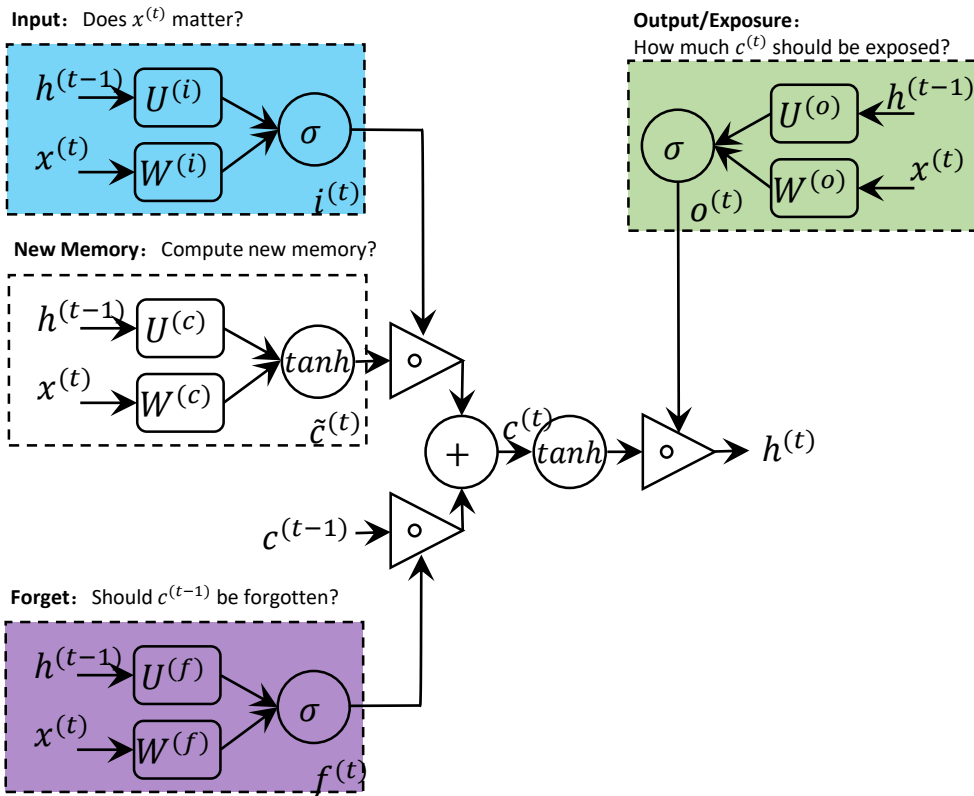


$$\delta_{in,1} = \delta_{out,t} \times \frac{\partial h_{out,t}}{\partial h_{in,t}} \times \frac{\partial h_{in,t}}{\partial h_{out,t-1}} \times \dots \times \frac{\partial h_{in,2}}{\partial h_{out,1}} \times \frac{\partial h_{out,1}}{\partial h_{in,1}}$$

Vanishing Gradient Problem



LSTM: Long Short Term Memory



$$i^{(t)} = \sigma(W^{(i)}x^{(t)} + U^{(i)}h^{(t-1)})$$


$$f^{(t)} = \sigma(W^{(f)}x^{(t)} + U^{(f)}h^{(t-1)})$$

$$o^{(t)} = \sigma(W^{(o)}x^{(t)} + U^{(o)}h^{(t-1)})$$

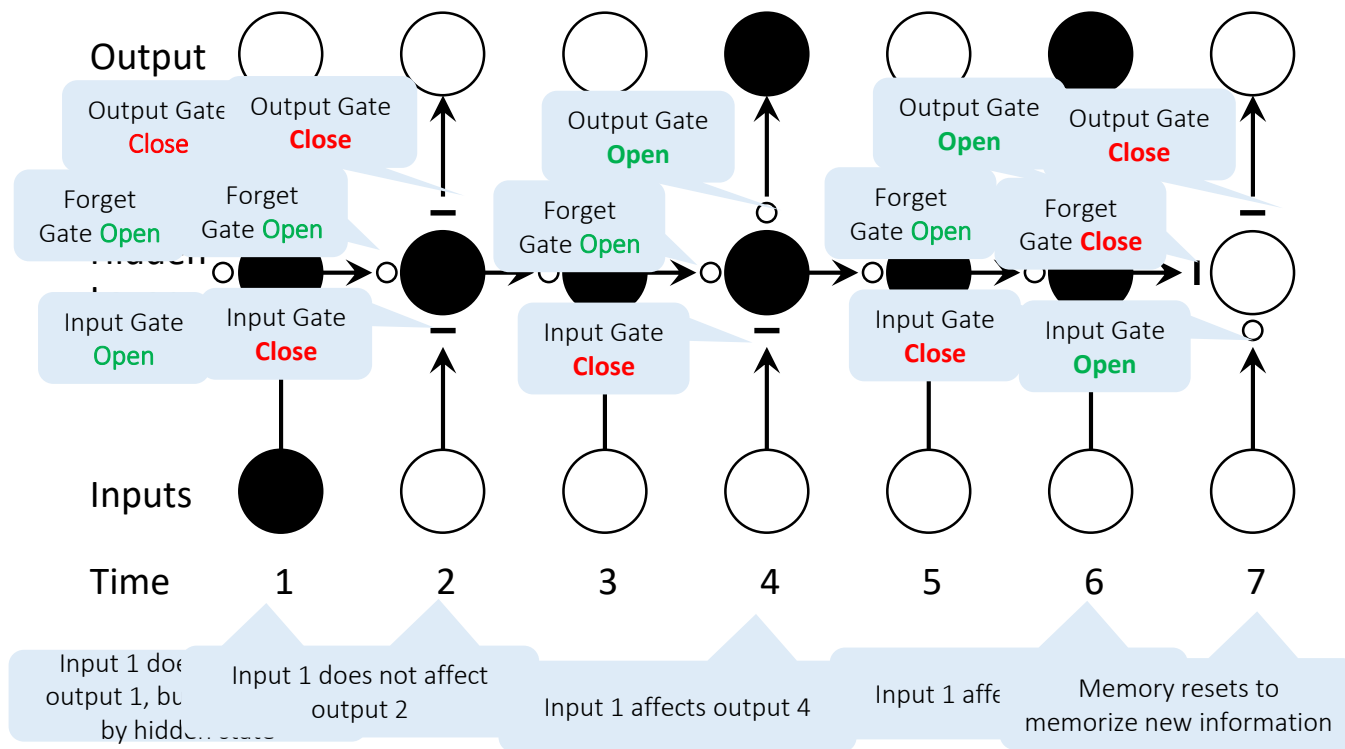
$$\tilde{c}^{(t)} = \tanh(W^{(c)}x^{(t)} + U^{(c)}h^{(t-1)})$$

$$c^{(t)} = f^{(t)} \circ \tilde{c}^{(t-1)} + i^{(t)} \circ \tilde{c}^{(t)}$$


$$h^{(t)} = o^{(t)} \circ \tanh(c^{(t)})$$



LSTM: Long Short Term Memory



LSTM: An Example



Feed Forward vs. LSTM LMs

- **Task:** Train FFNN and LSTM LMs on same data
- **Result:** LSTM outperforms FFNN, even if FFNN uses long context

<i>n</i> -gram Order	Num Hidden Layers	Perplexity
5	1	65.8
5	3	58.9
7	3	55.2
10	3	52.8
15	3	51.9
20	3	51.6
LSTM	1	45.1
LSTM	2	41.8

English LM, 100M words, 10k output vocab

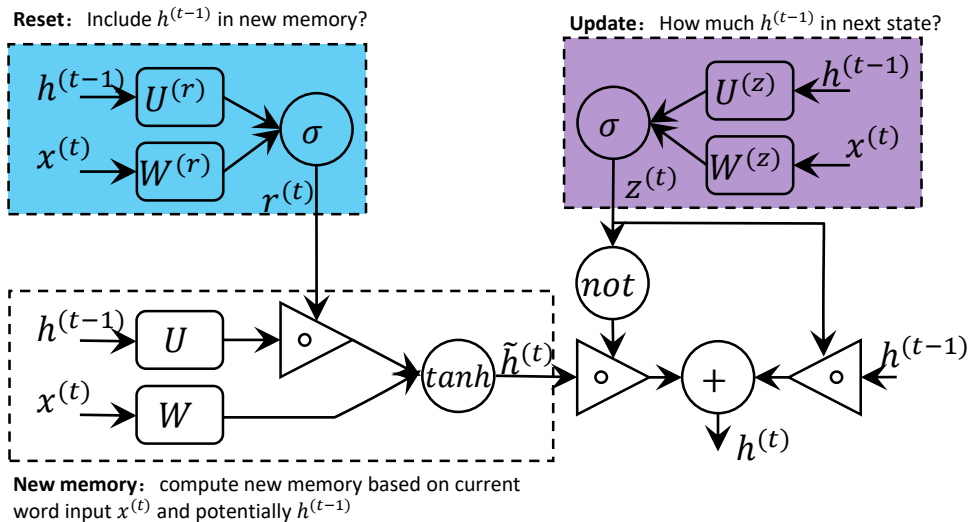
Feed Forward vs. LSTM LMs

- Qualitative analysis: LSTM is much better at “parsing” the input

Segment	20-gram FF Log Prob	Recurrent Log Prob
the lawsuit , filed wednesday on behalf of linda and robert lott of birmingham , alleges	-8.9	-2.7
the lawsuit alleges	-2.9	-2.8
some journalists said the claim that instant news was more incendiary than reports delivered more slowly was	-9.3	-1.5
some journalists said the claim was	-1.8	-0.8

Word being predicted is in **bold**

GRU: Gated Recurrent Unit




$$z^{(t)} = \sigma(W^{(z)}x^{(t)} + U^{(z)}h^{(t-1)})$$

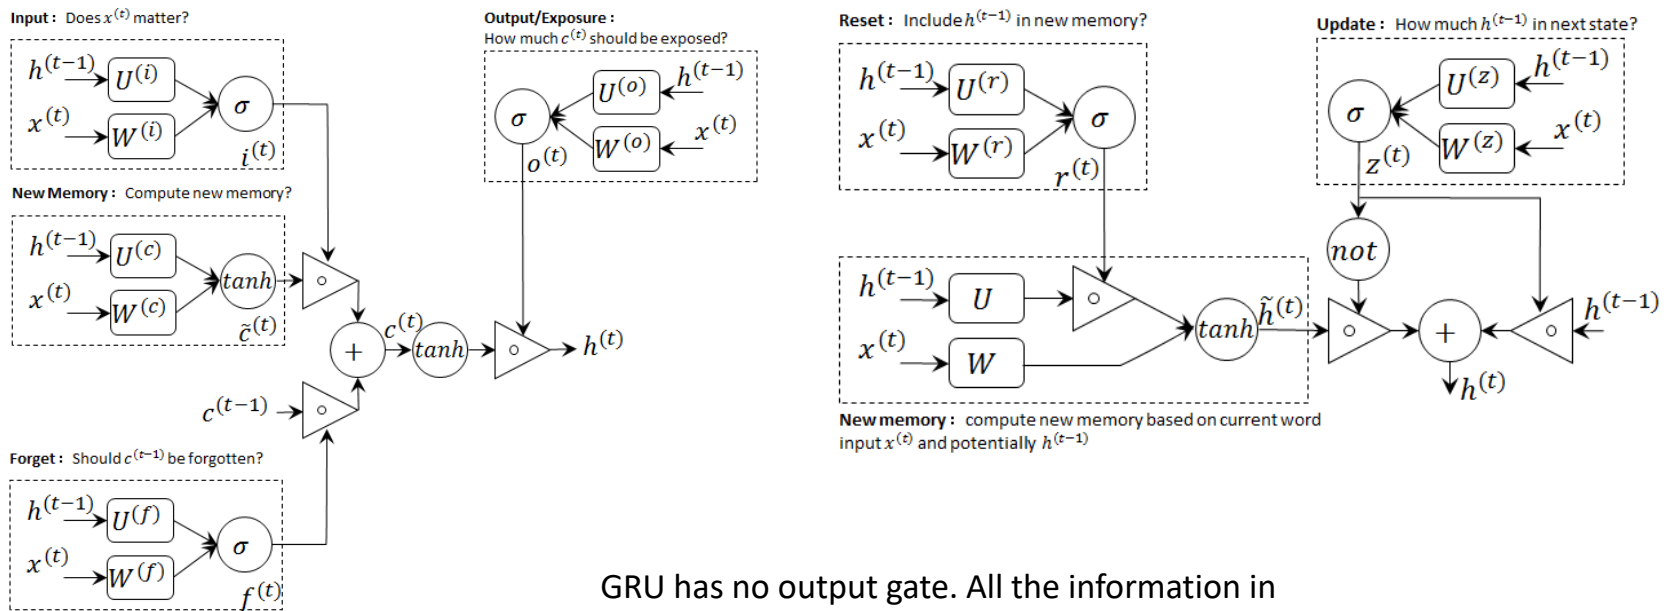
$$r^{(t)} = \sigma(W^{(r)}x^{(t)} + U^{(r)}h^{(t-1)})$$

$$\tilde{h}^{(t)} = \tanh(r^{(t)} \circ U h^{(t-1)} + W x^{(t)})$$

$$h^{(t)} = (1 - z^{(t)}) \circ \tilde{h}^{(t)} + z^{(t)} \circ h^{(t-1)}$$



LSTM vs GRU




GRU has no output gate. All the information in memory/hidden state will affect the output. The only way to make information inactive is to reset/forget it.


Outline

- Representation Learning First
- Introduction to DNN
- DNN for Natural Language Processing
 - DNN for Word Embedding
 - DNN for Language Modeling
 - **DNN for Machine Translation**

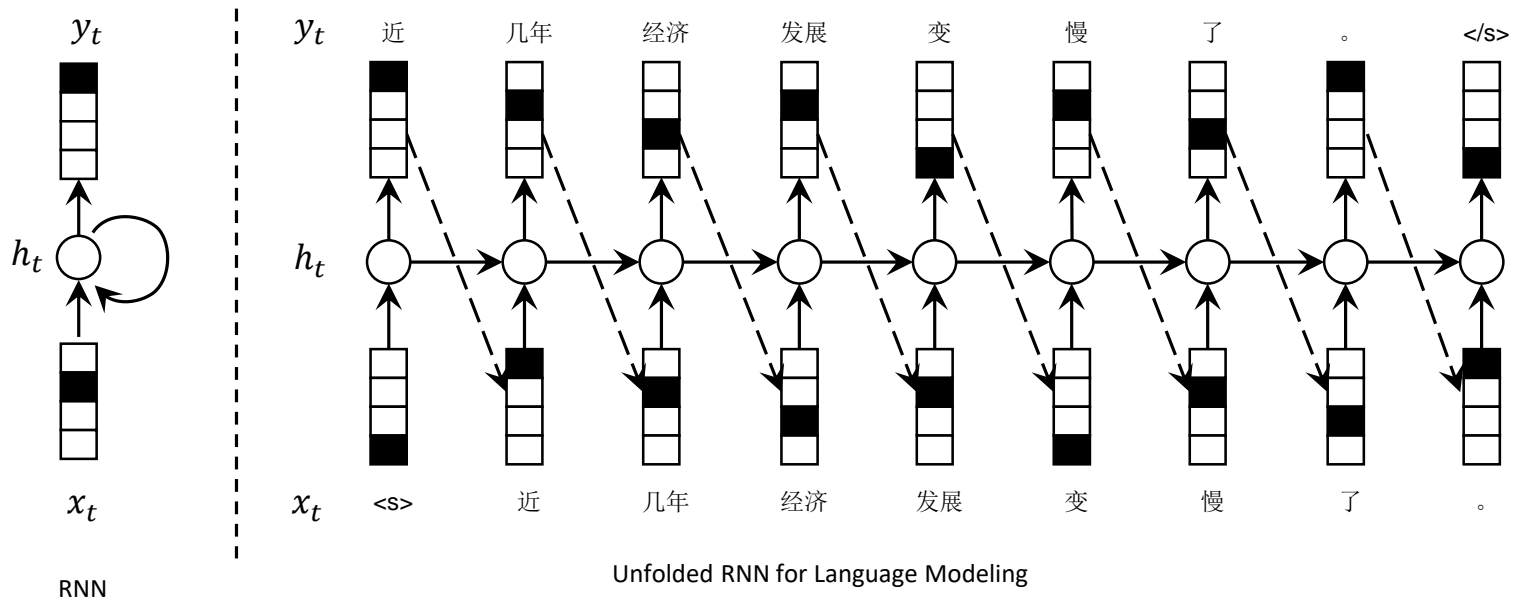
Encoder-Decoder for NMT




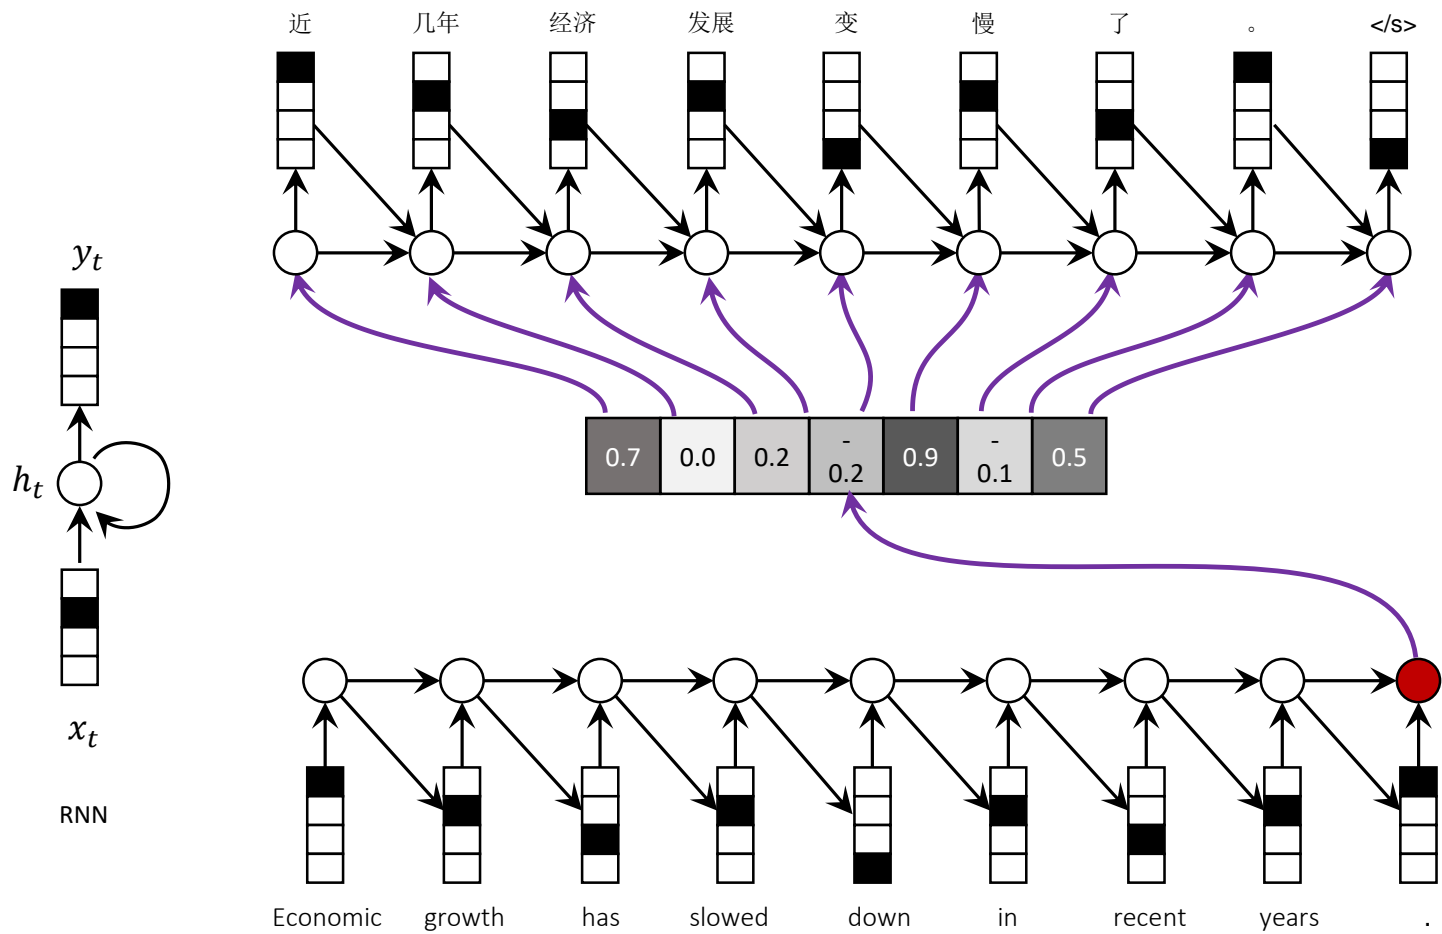
RNN-based Encoder-Decoder for NMT




RNN: Recurrent Neural Network










Encoder-Decoder for NMT





Motivation of Attention

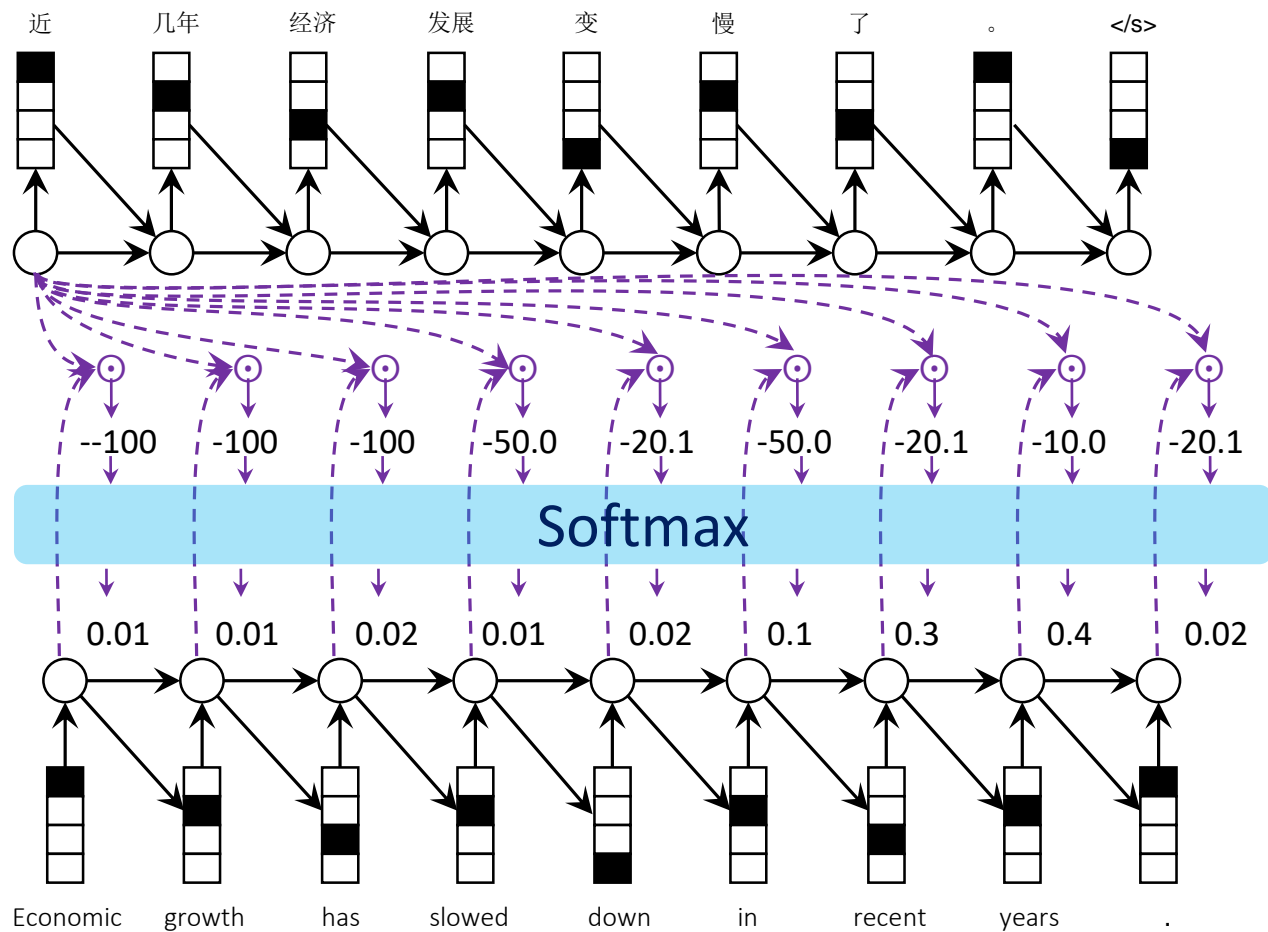
- NMT performs bad for long sentences
 - A long way from source to target
 - Only last hidden vector for decoding
 - Fixed-length hidden state is not enough
- Solution
 - Connect the source and target directly
 - Use all the hidden states for decoding

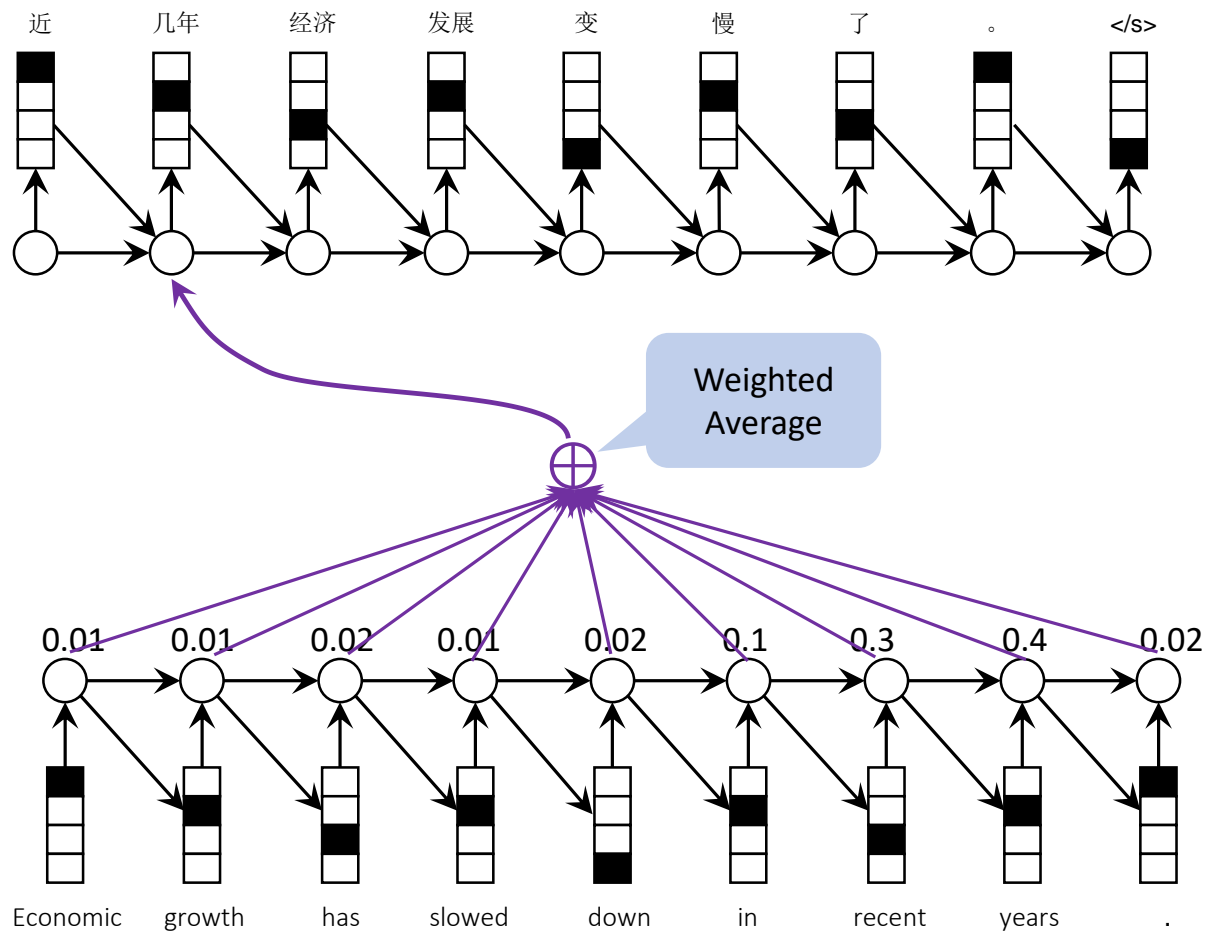















Attention based Encoder-Decoder



Attention based Encoder-Decoder

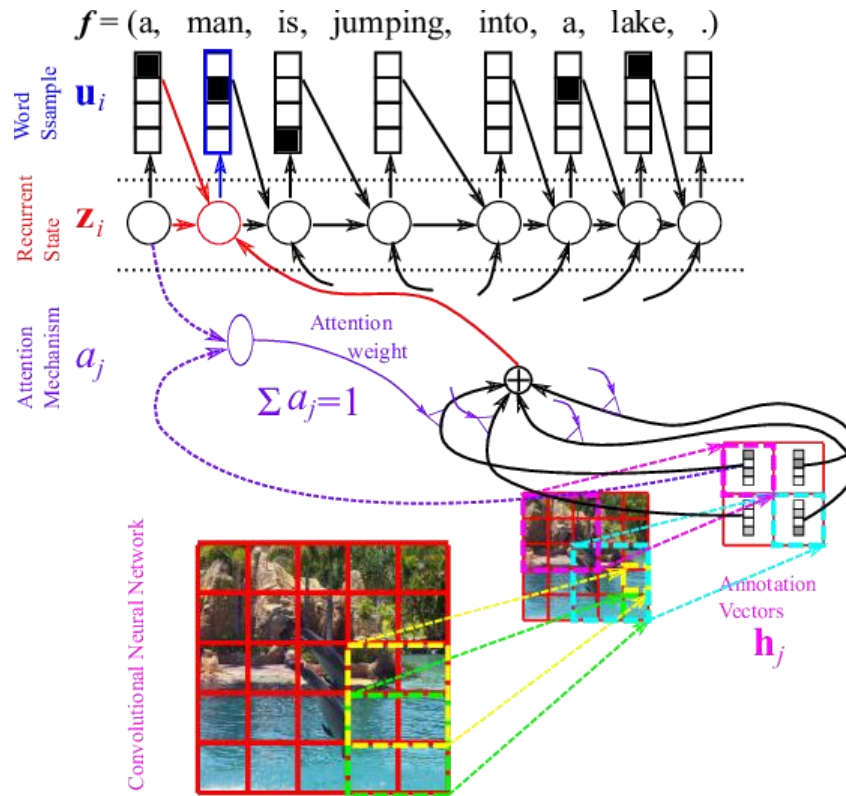


Case Study of Attention



$$\alpha_{ij} = \frac{\exp(e_{ij})}{\sum_k \exp(e_{ik})}$$

Image Caption Generation with Attention



not learn good alignments: the *global (location)* model can only obtain a small gain when performing unknown word replacement compared to using other alignment functions.¹⁴ For *content-based* functions, our implementation *concat* does not yield good performances and more analysis should be done to understand the reason.¹⁵ It is interesting to observe that *dot* works well for the global attention and *general* is better for the local attention. Among the different models, the local attention model with predictive alignments (*local-p*) is best, both in terms of perplexities and BLEU.

Case Study of Attention



A dog is standing on a hardwood floor.



A stop sign is on a road with a mountain in the background.



A group of people sitting on a boat in the water.



A giraffe standing in a forest with trees in the background.

$$\alpha_{ij} = \frac{\exp(e_{ij})}{\sum_k e_{ik}}$$

Thanks